Aurora OS

0. Product Name & One-Liner

(Working name; swap later.)

Name: Aurora Research OS
One-liner:

> AI-assisted, template-locked “Research-in-a-Box” that turns a clinician’s idea into a coherent, standards-aligned **baseline draft package** and live eCRF workspace—so good ideas don’t die before Ethics.

---

1. First Principles (Everyone Agrees)

P1. Clinician-first.
Start from how clinicians naturally talk/think. No stats jargon gatekeeping. The system translates clinical reasoning into structured research.

P2. Baseline, not black box.
We generate **strong drafts**, not decrees. Everything is editable. PI + IEC remain responsible.

P3. Constrained, not “anything goes”.
v1 only supports well-understood designs:

- Prospective / retrospective cohort
- Simple registries
- Simple 2-arm RCTs
- Diagnostic accuracy studies

Out-of-scope ideas are explicitly called out as such.

P4. Transparent & auditable.
Every assumption, formula, and AI decision traceable. No magic.

P5. Ethics & safety by design.
We enforce presence of core ethical elements and block obviously unsafe / nonsensical setups.

P6. Plug-in humans, don’t replace them.
Easy share flows for statisticians, IEC members, research offices to review and improve the

draft.

This is the lens through which all specs below are set.

---

2. Core Value Proposition

For clinicians:
Turn, “We should study this,” into a **concrete starting package in hours**, without leaving your clinical comfort zone.

For IECs & biostatisticians:
Receive coherent, structured proposals with transparent assumptions instead of chaotic Word docs.

For institutions:
Standardized, auditable study workflows that lower friction and raise baseline quality.

---

3. In-Scope vs Out-of-Scope (Hard Lines)

In v1 (must work well):

- Supported study types:
  - Cohort (prospective/retrospective)
  - Disease/quality registry
  - Simple 2-arm randomized trial
  - Diagnostic accuracy
- Adult, non–high-risk scenarios as default.
- Drafting:
  - Protocol
  - Sample size justification
  - SAP
  - CRF / eCRF schema
  - PIS/ICF
  - IEC cover note
  - Compliance checklist
  - Registry mapping sheet (e.g., CTRI-style fields)
- Live:
  - Study workspace
  - eCRFs + roles
  - Audit trail

- Safe exports

Out of v1:

- Adaptive/Bayesian / complex multi-arm trials.
- High-risk first-in-human / gene therapy style stuff.
- Automatic CTRI/DCGI/IRB submission.
- Any claim of legal / regulatory approval.

UI must clearly say when an idea is **not supported** by v1; no fake confidence.

---

4. Functional Spec (Bible Version)

4.1 FR-0: Clinician Comfort Zone Intake

- Intake is **plain language**:
  - “What problem are you seeing?”
  - “In which patients?”
  - “What do you want to measure/change?”
  - “Over what timeframe?”
  - “What’s realistic in your setting?”
- System converts this to structured JSON (`StudyDefinition`):
  - Condition, population, setting
  - Candidate primary outcome
  - Candidate design from supported set
- Always show:
  - A human-readable “Study Story” alongside structured fields for confirmation.

4.2 FR-1: Design Selection & Validation

From `StudyDefinition` , system:

1. Suggests **one or more supported design patterns**.
2. For chosen pattern, proposes:
   - Primary/secondary objectives
   - Endpoint definitions
   - Inclusion/exclusion criteria
   - Follow-up schedule
   - Core variable list

Constraints:

- All options come from curated **Template Libraries**.
- User can edit but within validation:
  - Every primary endpoint must map to:
    - A variable in the CRF
    - A planned analysis in the SAP

If broken → blocking warnings until resolved/acknowledged.

4.3 FR-2: Deterministic Sample Size Engine

- Implemented purely in backend code (no LLM math).
- Supports:
  - Single-arm precision
  - Two-arm comparisons (means/proportions)
  - Basic survival (log-rank)
  - Diagnostic accuracy (Se/Sp, etc.)

Behaviour:

- Ask user for assumptions or offer **clearly labeled suggestions**.
- Output:
  - N per group/overall
  - Scenario table
  - Auto-generated explanation text sourced from stored inputs.
- Store:
  - Formula choice
  - Parameters
  - Version

LLM may explain; **cannot** alter the numbers.

4.4 FR-3: Document Factory (Baseline Draft Package)

On command (“Generate Baseline Package”), system produces **drafts**:

1. **Protocol**
   - Standard headings (Background, Rationale, Objectives, Methods, Safety, Data Mgmt, etc.).

- Populated from structured spec + template snippets.
- Mandatory sections **not deletable** (can be edited, not removed).

2. **SAP**
   - Derived from:
     - Design type
     - Endpoint types
     - Sample size rationale
   - Outlines:
     - Analysis populations
     - Primary/secondary analysis methods
     - Handling of missing data (from templates).
3. **CRF Pack**
   - Auto-build from variable list:
     - Screening, Baseline, Procedure/Exposure, Follow-ups, Outcomes.
   - Exports:
     - Printable PDFs
     - JSON schema for eCRFs.
4. **PIS/ICF**
   - Templates containing all essential elements.
   - Auto-filled with study-specific info.
   - Ready for localization.
5. **IEC Cover Note**
   - Short study summary, design overview, risk/benefit.
6. **Compliance Checklist**
   - Presence (not legal sufficiency) of:
     - Objectives, endpoints, methods
     - Consent, risks, confidentiality, data protection, compensation, etc.
   - Traffic-light style.
7. **Registry Mapping Sheet**
   - Key study fields mapped to public registry-style fields.
   - Flags values PI must add manually (IEC number, DCGI status, etc.).

Mandatory:
Every generated doc has:

- Footer / cover text:
  - “Draft generated by Aurora Research OS based on user inputs.”
  - “Requires review and approval by Principal Investigator and Ethics Committee.”
- No “approved/ certified/ guaranteed” language.

4.5 FR-4: Compliance Gate (“No Blind Launch”)

Before a study can be “Locked & Launched”:

- System runs compliance & consistency checks:
  - Objectives ↔ Outcomes ↔ SAP aligned?
  - PIS/ICF essential clauses present?
  - Study type supported?
- Critical issues:
  - Must be resolved OR explicitly acknowledged via checkbox by PI.
- Only then:
  - Study Version can be locked.
  - Live workspace is generated.

This “gate” is non-negotiable.

4.6 FR-5: Risk & Sanity Guardrails

- If user proposes clearly unethical/unsafe/unsound structure:
  - E.g., high-risk intervention, no safety oversight, no consent.
- System must:
  - Flag with strong warning.
  - Stop auto-normalizing.
  - Suggest modifications or “Expert review required”.
- Logs these events in audit trail.

4.7 FR-6: Study Workspace & eCRFs

When locked:

- System creates **Study Workspace**:
  - Configured eCRFs from CRF schema.
  - Roles:
    - Owner (PI)
    - Editor (Co-I)
    - Data Entry
    - Monitor / Read-only
- Features:
  - Add / edit participant records.

        - Track status: screened / enrolled / completed.
        - Basic dashboard:
            - Enrollment
            - Follow-up
            - Data completeness
- Export:
    - CSV + codebook
    - JSON schema
    - Audit report PDF (config + changes).

Any protocol amendment:

- Creates new StudyVersion.
- Diffs logged (IEC amendment-ready).

4.8 FR-7: Collaboration & Handoff

- “Share for Review” link:
    - Grants read/comment access to:
        - Statisticians
        - IEC/research cell
    - Comments are tracked; no silent overwrites.
- Final control:
    - PI confirms final content.
    - All changes versioned.

This operationalizes “baseline they can build off”.

---

5. Non-Functional Spec

- **Security**
    - HTTPS, TLS
    - Encryption at rest
    - Role-based access
    - Study-level isolation
- **Data Residency**
    - Region selection (e.g., India region for Indian clients).
- **Auditability**
    - Immutable event log:

        - Who changed what, when.
        - Which AI/template version used.
- **Performance**
    - Baseline package generation typically < 60s.
    - eCRF interactions fast (< 300ms typical).
- **Reliability**
    - 99.5% uptime target for core app.
- **No PHI in model training**
    - Production runs on non-retaining LLM endpoints or self-hosted models for PHI zones.

---

6. AI / ML Architecture (Guardrailed)

Components:

1. **Idea Parser**
    - LLM → structured PICO/spec JSON.
    - Validated by schema.
2. **Design Assistant**
    - Runs against **whitelisted templates**.
    - Can only pick/edit inside allowed schema.
3. **Document Composer**
    - Uses structured spec + snippet libraries.
    - No freeform regulatory inventions.
4. **Compliance Checker**
    - LLM + rules:
        - Only flags gaps/inconsistencies.
        - Does NOT claim compliance.

Guardrails:

- Low temperature for structural outputs.
- JSON schema enforcement.
- Deterministic stats outside LLM.
- Every AI step recorded in `AuditEvent`.

---

7. System Architecture (High-Level)

- **Frontend:**
  - Next.js / React
  - Tailwind + component library (e.g., shadcn-style)
- **Backend:**
  - Typed framework (NestJS / similar)
  - Services:
    - Auth / Org
    - Study & Templates
    - Orchestrator (AI pipeline)
    - Docs (DOCX/PDF)
    - CRF & Data
    - Audit Log
- **Data:**
  - Postgres for core entities
  - S3-compatible storage for docs
  - Redis for cache/queues

---

8. Canonical Data Entities

- `User`
- `Organization`
- `Study`
- `StudyVersion`
- `Template` (designs, CRFs, PIS/ICF blocks)
- `Document` (type, version, URL)
- `CRFSchema`
- `ParticipantRecord`
- `AuditEvent` (actor, action, entity, timestamp, payload)

This is the minimum model; extend carefully.

---

9. Golden Path (Demo & Real Use)

1. Clinician: types natural-language idea.
2. System: maps to supported design; clinician confirms.
3. System: guides through endpoints, follow-up, core variables.
4. Sample size: runs deterministic calc, shows explanation.

5. “Generate Baseline Package”:
    - Protocol, SAP, CRFs, PIS/ICF, IEC note, checklist, mapping sheet.
6. Compliance gate: clinician reviews, fixes/acknowledges.
7. “Lock & Launch”:
    - Study workspace + eCRFs go live.
8. Same day:
    - Enter first patient.
    - Export full draft pack + audit summary → send to IEC.

That’s your promise. That’s what everyone on the (real or virtual) team is building.